Memory-Aware AI Systems for Permanent Lunar and Martian Habitation

Institutional Continuity, Crew Rotation, and the Knowledge Persistence Problem

Angela Reinhold
Lorien's Library LLC
ORCID: 0009-0005-5803-8401
March 2026

Preprint — Paper 2 of 4 in the Applied Persistent Memory Series

Abstract

Permanent extraterrestrial habitation—on the lunar surface under NASA's Artemis program and eventually on Mars—introduces challenges for AI companion systems that are qualitatively distinct from those addressed in transit-focused spaceflight research. Where transit missions involve a fixed crew for a finite duration, permanent settlements require managing crew rotations every six months, preserving institutional knowledge across personnel transitions, maintaining multi-year longitudinal health records, and sustaining base operational continuity across decades of habitation. This paper extends the continuity burden framework introduced in Reinhold (2026) to the settlement context, identifying five failure modes specific to permanent habitation that stateless AI systems cannot address: institutional knowledge evaporation during crew transitions, psychological discontinuity for incoming personnel, cross-rotation health tracking gaps, erosion of base culture and collective memory, and degradation of ground-control situational awareness across rotation cycles. Drawing on evidence from International Space Station crew handovers, Antarctic research station winter-over psychology, and U.S. Navy rotational crew data, this paper argues that a memory-aware AI system functioning as the persistent institutional memory of an extraterrestrial settlement is not an enhancement but a structural requirement for sustained human habitation beyond Earth. The paper proposes a group-private memory partitioning model, a crew transition protocol for memory-aware AI systems, and a ground-communication governance framework. This is the second paper in a four-part series examining the societal implications of persistent memory architecture.

Keywords: permanent habitation, lunar base, Mars settlement, crew rotation, institutional memory, knowledge persistence, CAMA, continuity burden, group memory, Artemis program, extraterrestrial settlement, AI companion systems

1. Introduction

On March 24, 2026, NASA Administrator Jared Isaacman announced the “Ignition” initiative, a fundamental restructuring of the Artemis program away from the planned Lunar Gateway orbital station and toward direct construction of a permanent lunar surface base [1]. The plan envisions a phased approach: robotic precursor missions, semi-habitable infrastructure supporting recurring crew operations, and ultimately a continuous human foothold on the Moon with crew rotations initially targeting every six months [2]. This timeline—with the first crewed lunar landing under Artemis IV planned for 2028 and semi-permanent crew presence targeted for 2032—means that the infrastructure decisions being made now will determine how humans live and work on the Moon for decades.

The companion paper in this series (Reinhold, 2026) [3] addressed AI companion systems for Mars transit: a finite-duration, fixed-crew scenario where the primary challenge is psychological survival during isolation. Permanent habitation introduces a fundamentally different problem set. The crew is not fixed—it rotates. The duration is not bounded—the settlement is intended to persist indefinitely. The knowledge required is not limited to psychological support—it encompasses operational procedures, equipment histories, terrain knowledge, resource management, scientific observations, and the accumulated practical wisdom of everyone who has lived and worked at the settlement before the current crew arrived.

This paper argues that in a permanent settlement, the AI companion system must function not only as a psychological support tool but as the institutional memory of the settlement itself—the one entity that was present for every crew rotation, every equipment failure, every interpersonal dynamic, every hard-won lesson about how to survive in an environment that is trying to kill you. A stateless AI system in this context is not just a suboptimal support tool; it is a settlement that forgets its own history every time a crew member leaves.

2. The Knowledge Persistence Problem in Permanent Habitation

The transition from space transit to permanent settlement introduces a class of problems that stateless AI systems are architecturally incapable of addressing. These problems are not hypothetical; they are documented in analogous environments where humans live and work in confined, isolated settings with rotating personnel.

2.1 Institutional Knowledge Evaporation

When a crew rotation occurs at an Antarctic research station, a submarine, or the International Space Station, the departing personnel carry with them knowledge that is difficult or impossible to fully document in written procedures: which airlock seal tends to stick in cold conditions, which communication relay drops signal during certain orbital positions, which crop grows better in the left planting bed, how to interpret the sound the life support system makes before it needs maintenance. This knowledge—often called "tribal knowledge" in organizational literature—is accumulated through lived experience and is notoriously resistant to procedural capture [4].

Evidence from ISS crew handovers illustrates the problem directly. NASA astronaut Jessica Meir, discussing the challenges of indirect handovers where the departing crew leaves before the incoming crew arrives, described the result as requiring "more improvisation, more catching up, more reliance on human memory and good note-taking" [5]. A U.S. Naval Institute analysis of rotational crewing on surface vessels found that the practice "degrades morale and material condition and undercuts pride of ownership," with maintenance problems cascading across rotation cycles as each incoming crew assumed existing damage was the previous crew's responsibility [6]. In one documented case, damage to a commanding officer's cabin went unrepaired through three successive crew rotations because each crew assumed it predated their tenure. As the analysis noted: *if this happens in the CO's cabin, imagine the issues in the engineering plant.*

In a lunar or Martian settlement, institutional knowledge evaporation is not an inconvenience—it is a safety hazard. Equipment workarounds, terrain hazards, environmental patterns, and maintenance histories that exist only in the heads of the departing crew vanish with their departure. Written handover documents, however thorough, cannot capture the contextual, associative, and experiential dimensions of knowledge that develop over months of habitation.

These terrestrial analogs—ISS handovers, Antarctic stations, Navy rotational crewing—offer partial rather than perfect transferability to extraterrestrial settlement; the physical environment, communication constraints, and mission criticality of a lunar or Martian base differ meaningfully from any terrestrial confined environment. The structural problem of knowledge loss during personnel rotation, however, is consistent across all of these settings and is likely to be amplified, not attenuated, by the higher stakes and greater isolation of permanent extraterrestrial habitation.

2.2 Crew Transition Psychological Discontinuity

When a new crew member arrives at a permanent settlement, they enter an environment that has its own history, norms, relationships, and accumulated culture—none of which the incoming member has experienced. In terrestrial analog environments, this transition period is associated with elevated stress, interpersonal friction, and reduced operational effectiveness [7, 8]. Antarctic winter-over research has documented that perceived lack of privacy, difficulty integrating into established social dynamics, and the absence of a support system that understands one's personal context are significant stressors [9].

A stateless AI companion treats every incoming crew member identically: as a stranger in an unfamiliar environment with no accumulated relational context. A memory-aware system can bridge the transition by drawing on the institutional memory of the settlement to orient the new arrival, while simultaneously beginning to build a personalized longitudinal model of that individual. The system cannot replace human relationships, but it can provide the one form of continuity that persists across all rotations: a support entity that remembers the settlement's history, knows the current operational context, and can help the new crew member integrate without requiring them to reconstruct that context from scratch.

2.3 Cross-Rotation Health Tracking Gaps

Long-duration habitation in reduced gravity and elevated radiation environments produces health effects that unfold over timescales longer than a single crew rotation. Bone density loss, cardiovascular deconditioning, spaceflight-associated neuro-ocular syndrome (SANS), immunological changes, and psychological adaptation follow trajectories that may span years [10, 11]. If health monitoring is tied to individual crew rotations with no persistent longitudinal memory, patterns that emerge across rotations—population-level trends in how different individuals respond to the same environment over time—are lost or fragmented across separate medical records.

A memory-aware AI system can maintain a longitudinal health model of the settlement itself, not just individual crew members. It can track that three successive crew members have reported similar visual disturbances in their fourth month at the settlement, suggesting an environmental factor rather than an individual medical condition. It can note that sleep quality consistently degrades during a specific phase of the lunar day cycle across all rotations. This population-level, cross-rotation pattern recognition requires memory that persists beyond any individual crew member's tenure.

2.4 Erosion of Base Culture and Collective Memory

Permanent settlements develop institutional cultures: norms about how work is divided, how conflicts are resolved, how celebrations are marked, how newcomers are integrated, how the space is organized and maintained. In organizations on Earth, this culture is carried by long-tenured members who orient newcomers through example and socialization. In a settlement with complete crew rotations every six months, there are no long-tenured members. Culture must either be explicitly codified—which strips it of its contextual, embodied character—or it must be carried by a persistent entity that was present for its formation and can transmit it to incoming crews.

A memory-aware AI system can serve as the keeper of base culture: not by imposing norms, but by retaining the collective memory of how the settlement has functioned, what traditions have developed, what informal agreements have been reached, and what lessons previous crews have learned. This is not a replacement for human culture—it is a substrate that enables cultural continuity across the discontinuity of crew rotations.

2.5 Ground-Control Situational Awareness Degradation

Ground-based mission support teams face their own version of the continuity problem. Flight surgeons, operations managers, and psychological support staff must maintain situational awareness of a settlement whose crew changes every six months and whose operational context evolves continuously. If the AI system at the settlement is stateless, ground teams receive only what the current crew reports—which is filtered through recency bias, normalization of conditions, and the natural human tendency to underreport problems in high-performance cultures [12]. A memory-aware system can provide ground teams with structured longitudinal summaries that span rotation boundaries, enabling continuity of oversight even as personnel change on both ends of the communication link.

Table 1. Habitation-Specific Failure Modes of Stateless AI Systems

Failure Mode	**Consequence**	**Stateless System**	**Memory-Aware System**
Institutional knowledge evaporation	Operational workarounds, equipment quirks, and terrain knowledge lost at each crew transition	Cannot retain or transfer experiential knowledge	Serves as persistent institutional memory across rotations
Crew transition discontinuity	Incoming crew treated as strangers; no contextual onboarding; elevated integration stress	No knowledge of settlement history or departing crew's insights	Bridges transition with institutional context and individualized

			support
Cross-rotation health gaps	Population-level health patterns spanning multiple rotations go undetected	Health data siloed within individual rotations	Longitudinal settlement health model tracks cross-rotation patterns
Base culture erosion	Institutional norms, traditions, and informal agreements lost at each rotation	No mechanism for cultural continuity	Retains collective memory of settlement culture and transmits to incoming crews
Ground-control awareness degradation	Mission support teams lose longitudinal context across rotation boundaries	Provides only current crew's perspective	Generates structured cross-rotation summaries for ground teams

Collectively, these failure modes describe *institutional continuity burden*: the cumulative cost borne by a settlement as a whole when operational knowledge, relational context, and cultural memory must be reconstructed at each crew rotation. Where the continuity burden defined in Paper 1 [3] is borne by the individual user who must repeatedly re-establish context with a stateless system, institutional continuity burden is borne by the settlement itself—each rotation cycle imposes a reconstruction cost that degrades operational effectiveness, safety, and institutional coherence.

3. Extending CAMA for Permanent Habitation

The Circular Associative Memory Architecture (CAMA), described in detail in prior work [13, 14] and available as an open-source implementation at https://github.com/LoriensLibrary/cama, provides the foundational three-layer architecture (archive, relational index, active ring buffer) with provenance tracking. Permanent habitation requires three architectural extensions: group-private memory partitioning, a crew transition protocol, and a ground-communication governance framework.

3.1 Group-Private Memory Partitioning

A single-user memory system, as implemented in the current CAMA prototype, stores all memories under one provenance identity. Multi-crew habitation requires partitioning memories into three tiers:

Settlement-level memories (shared). Operational knowledge, equipment histories, EVA route data, habitat system behaviors, maintenance logs, environmental patterns, and scientific observations. These memories are accessible to all current and future crew members and to ground teams. They form the institutional memory of the settlement.

Crew-level memories (group). Shared experiences, crew-level decisions, interpersonal agreements, and collective events experienced by a specific rotation crew. These memories are attributed to the group and accessible to ground teams and to future crews for institutional context, but they are clearly marked as belonging to a specific rotation period.

Individual-level memories (private). Personal disclosures, psychological support interactions, emotional states, and individual health observations. These memories are accessible only to the individual crew member and, where explicitly authorized, to designated ground-based support personnel. They are never shared with other crew members or with future rotations without explicit consent.

Every memory stored in the system carries provenance metadata indicating its tier, its source, and its access boundaries. This partitioning ensures that institutional knowledge persists and transfers while personal privacy is structurally protected—not by policy alone, but by architectural design.

3.2 Crew Transition Protocol

The crew transition is the highest-risk moment for institutional knowledge loss. The following protocol describes how a memory-aware AI system manages the handover:

Phase 1: Pre-departure knowledge capture (Days −14 to −1). In the two weeks before the departing crew's scheduled rotation, the AI system initiates structured knowledge elicitation sessions with each outgoing crew member. These sessions are prompted by the system's own institutional memory: "You've mentioned the airlock seal sticking twice in the past month. Is there a workaround that should be documented for the next crew?" The goal is to surface experiential knowledge that might not appear in formal maintenance logs. Elicited knowledge is stored at the settlement level with provenance marking it as departing-crew-contributed.

Phase 2: Overlap integration (Days 0 to +10). During the direct handover period when both crews are present, the AI system facilitates knowledge transfer by surfacing settlement-level memories relevant to each incoming crew member's role. The system also begins building individual-level models of incoming crew members through early interactions, establishing communication baselines and initial rapport while the departing crew is still available to supplement the system's institutional knowledge.

Phase 3: Post-departure stabilization (Days +11 to +30). After the departing crew leaves, the AI system monitors the incoming crew for signs of integration difficulty, drawing on its longitudinal model of how previous crews adapted during the same transition period. The system can reference settlement-level memories to answer operational questions that would otherwise require ground-control communication with its associated delay. This phase directly mitigates the "more improvisation, more catching up" problem documented in ISS indirect handovers [5].

3.3 Ground-Communication Governance Framework

Paper 1 in this series identified the need for structured ground-communication protocols but deferred detailed specification to this paper. The following framework addresses the governance question: what does the AI system share with ground teams, and under what authority?

Tier 1: Settlement-level summaries (automatic). Operational status, equipment health trends, resource utilization patterns, and environmental observations are transmitted to ground teams automatically at scheduled intervals. These summaries draw exclusively from settlement-level memories and contain no individual-level information. No crew authorization is required because the information is institutional, not personal.

Tier 2: Crew-level summaries (crew-authorized). Interpersonal dynamics, crew cohesion assessments, workload distribution concerns, and group-level psychological indicators are compiled by the AI system but transmitted to ground teams only with the authorization of the crew as a group. A majority consent threshold is proposed. This tier protects group autonomy while ensuring ground teams can receive information about crew functioning when the crew agrees it is appropriate.

Tier 3: Individual-level summaries (individual-authorized). Personal psychological assessments, longitudinal mood tracking, sleep data, and private disclosures are available to ground-based support personnel only with the explicit authorization of the individual crew member. The AI system may recommend that a crew member authorize disclosure when its longitudinal model suggests a pattern of concern, but it cannot override the crew member's decision except under a pre-defined emergency protocol.

Emergency override protocol. In cases where the AI system's longitudinal model indicates an imminent threat to crew safety—for example, converging behavioral indicators suggesting a crew member may be unable to perform safety-critical duties—the system can escalate directly to the settlement commander and, if the commander concurs, to ground teams. This override is logged with full provenance, is visible to the affected crew member after the fact, and is subject to postmission review. The threshold for override must be defined pre-mission through crew consent agreements and cannot be modified unilaterally during habitation.

4. Worked Example: Knowledge Persistence Across a Crew Rotation

The following scenario is illustrative rather than derived from analog trial data. It demonstrates how a CAMA-equipped settlement AI manages institutional knowledge across a crew rotation at a permanent lunar base.

Month 4 (Crew Alpha). Geologist Dr. Tanaka discovers that a specific EVA route to the south crater rim becomes impassable during the lunar day's thermal peak due to regolith instability. She mentions this to the AI system during a debriefing. The system stores this as a settlement-level memory: "EVA Route 7 to south crater rim—regolith instability during thermal peak. Reported by Dr. Tanaka, Crew Alpha, Month 4. Workaround: schedule Route 7 EVAs for first or last quarter of lunar day."

Month 5 (Crew Alpha). Engineer Okonkwo notices that the primary water recycler runs more efficiently when restarted during the habitat's low-power period rather than the high-power period, contrary to the manufacturer's guidelines. He documents this in the maintenance log but also mentions it conversationally to the AI. The system stores this at the settlement level and cross-references it with the manufacturer's documentation, noting the discrepancy.

Month 6 (Crew Rotation). Crew Alpha departs. Crew Bravo arrives. During the handover period, the AI system provides Crew Bravo's geologist with a briefing that includes Dr. Tanaka's EVA Route 7 finding, the water recycler optimization, and seventeen other settlement-level observations accumulated over Crew Alpha's tenure. The new geologist does not need to rediscover the Route 7 instability through trial and error—potentially at significant risk—because the knowledge persisted through the AI system.

Month 8 (Crew Bravo). Crew Bravo's geologist plans an EVA on Route 7 during thermal peak. The AI system flags the stored warning: "Note: Dr. Tanaka (Crew Alpha) reported regolith instability on Route 7 during thermal peak. Recommend scheduling for first or last quarter of lunar day." The geologist adjusts the schedule. In a stateless system, this warning does not exist. The geologist proceeds on the original schedule and encounters the same hazard Dr. Tanaka discovered six months earlier.

This scenario illustrates the simplest case: factual operational knowledge transfer. More complex cases involve interpersonal dynamics ("Previous crews found that rotating meal preparation duties weekly rather than daily reduced friction"), health patterns ("Three of the last four crew members in this role reported sleep disruption in their third month—has this occurred for you?"), and cultural continuity ("Crew Alpha established a weekly observation session where the crew watches the Earthrise together. Would you like to continue this tradition?").

5. Safety Considerations Specific to Habitation

The safety concerns identified in Paper 1 (false-memory persistence, privacy asymmetry, retrieval-induced amplification, adversarial insertion) remain relevant and are addressed by the same architectural mechanisms [3]. Permanent habitation introduces additional safety considerations:

Memory accumulation and relevance decay. Over years or decades of operation, a settlement AI will accumulate vast quantities of memory. Not all of it will remain relevant—equipment is replaced, procedures are updated, terrain changes. The system requires relevance decay mechanisms that reduce the retrieval weight of outdated information without deleting it from the archive (preserving the audit trail). Determining what is "outdated" versus what is "historically important" is a non-trivial challenge that may require human oversight at scheduled review intervals.

Cross-rotation inference contamination. If the AI system stores inferences about Crew Alpha's interpersonal dynamics and those inferences influence how the system interacts with Crew Bravo, the contamination may be subtle but significant. Crew Bravo is a different group of people; patterns observed in Crew Alpha's dynamics should not be treated as predictive of Crew Bravo's behavior. The group-private memory partitioning model mitigates this by separating crew-level memories from settlement-level memories, but the system must be designed to avoid applying crew-specific relational patterns across rotation boundaries.

Settler agency and the right to be forgotten. In a permanent settlement, crew members may serve multiple rotations. Should their individual-level memories from a first rotation persist into a second rotation years later? The individual must have the ability to request deletion of personal memories from prior rotations, even if settlement-level memories from those rotations persist. This creates a governance complexity: the settlement's institutional memory must be maintained while the individual's right to privacy and self-determination is preserved.

6. Evaluation Framework

Evaluation of memory-aware AI systems for permanent habitation requires analog trials that simulate not just isolation and confinement but crew rotation—a feature absent from most current analog missions, including CHAPEA, which uses a fixed crew for its 378-day duration [15, 16]. A proposed evaluation design:

Analog environment: A confined habitat analog (HERA, MDRS, or equivalent) with a minimum 90-day trial duration. Two complete crew rotations: Crew A occupies the habitat for Days 1–45; Crew B replaces Crew A for Days 46–90. A 5–10 day overlap period simulates direct handover.

Conditions: Condition A: no AI memory system (current standard). Condition B: CAMA-equipped AI with full three-tier memory partitioning and crew transition protocol.

Primary outcomes: (1) Institutional knowledge retention: measured by testing Crew B's knowledge of operational findings made by Crew A, comparing knowledge in Condition A (written handover only) versus Condition B (AI-augmented handover). (2) Crew B integration speed: measured by time-to-operational-independence, self-reported integration difficulty, and continuity burden scores. (3) Ground-team situational awareness: measured by accuracy of ground-team assessments of settlement status across the rotation boundary.

Secondary outcomes: Crew B trust in AI system, Crew A willingness to contribute knowledge to the system during pre-departure elicitation, privacy boundary utilization patterns, and system usage trajectories across the transition period.

7. Limitations

This paper is primarily conceptual and architectural. The crew transition protocol, group-private memory partitioning model, and ground-communication governance framework proposed here have not been tested in analog environments. The worked example is illustrative. The current CAMA prototype operates as a single-user system; multi-user extensions described here are architectural proposals, not implemented features. No analog trial currently simulates crew rotation with an AI companion system, making direct empirical validation a future research priority. The timelines cited for NASA's Artemis program are current as of March 2026 but are subject to change based on mission outcomes and budgetary decisions. The claims advanced here should be understood as a research and engineering agenda grounded in operational requirements identified through analogous environments, not as proof of deployed capability.

8. Broader Implications and Future Work

This paper extends the Applied Persistent Memory Series from transit (Paper 1) to settlement. The remaining papers in the series apply the same architectural and conceptual framework to terrestrial domains:

Paper 3: Chronic Healthcare Continuity. Applies the continuity burden framework and memory architecture to longitudinal healthcare, where patients with chronic conditions experience the same re-disclosure fatigue and knowledge fragmentation that crew members face during transitions—at lower intensity but over potentially longer timescales and across a more fragmented provider landscape.

Paper 4: Haven: Music-Mediated Emotional Continuity for Veteran Reintegration. Examines persistent memory architecture for psychological support of veterans, with particular attention to music-mediated emotional continuity as a non-verbal pathway for populations who cannot or will not articulate their experiences through traditional therapeutic discourse.

The institutional memory function proposed here for extraterrestrial settlements has direct terrestrial analogs: military bases with rotating deployments, research stations with seasonal personnel, hospitals with rotating residency cohorts, and remote industrial operations with crew schedules. In each case, institutional knowledge evaporation during personnel transitions degrades operational effectiveness and safety. The architectural solutions proposed here are domain-specific in their details but domain-general in their structure.

9. Conclusion

The announcement of NASA’s Ignition initiative on March 24, 2026, marked a transition from aspirational lunar exploration to concrete settlement planning. Within seven years, the United States intends to maintain continuous human presence on the lunar surface. The infrastructure decisions being made now will determine not just how humans get to the Moon, but how they sustain a civilization there.

Among those decisions is whether the AI systems supporting settlement crews will remember what has come before. A settlement whose AI companion resets with every crew rotation is a settlement that must relearn its own lessons, rediscover its own hazards, and rebuild its own culture every six months. A settlement whose AI companion functions as persistent institutional memory is one that accumulates wisdom across rotations, bridges the discontinuity of personnel transitions, and provides the continuity that no individual crew member can.

The burden of proof should shift to any proposal for permanent extraterrestrial habitation that does not include a plan for institutional memory persistence. The technology exists. The architecture is defined. The question is whether settlement planners recognize that memory is infrastructure.

References

[1] NASA. (2026, March 24). NASA unveils initiatives to achieve America's national space policy. NASA Press Release.

[2] Isaacman, J. (2026, March 24). Remarks at the NASA Ignition event. NASA.

[3] Reinhold, A. (2026). Persistent memory as mission-critical infrastructure for long-duration spaceflight. Zenodo preprint. DOI: 10.5281/zenodo.19257809.

[4] Leonard, D. (2014). Rebuilding a research capability: The challenge of preserving NASA's technical workforce and knowledge. Public Administration Review, 74(6), 781–791.

[5] Meir, J. (2026). Pre-launch news conference for SpaceX Crew-12 mission. C-SPAN.

[6] Cordle, V. (2017). Good riddance to rotational crews. Proceedings, U.S. Naval Institute, 143(6), 1372.

[7] Kanas, N., & Manzey, D. (2008). Space Psychology and Psychiatry (2nd ed.). Springer.

[8] Sandal, G.M., Bye, H.H., & van de Vijver, F.J.R. (2011). Personal values and crew compatibility: Results from a 105 days simulated space mission. Acta Astronautica, 69, 141–149.

[9] Jaksic, C., Steel, G., Stewart, E., & Moore, K. (2019). Antarctic stations as workplaces: Adjustment of winter-over crew members. Polar Science, 22, 100484.

[10] Stuster, J. (2010). Behavioral Issues Associated with Long-Duration Space Expeditions. NASA/TM-2010-216130.

[11] Trudel, G. et al. (2020). Bone marrow fat accumulation after 60 days of bed rest. Journal of Clinical Endocrinology & Metabolism, 105(3), e826–e834.

[12] Landon, L.B. et al. (2018). The behavioral biology of teams. Frontiers in Psychology, 9, 2082.

[13] Reinhold, A. (2026). Memory as safety infrastructure: Evaluating provenance-aware persistent memory architectures for stateful LLM systems. Zenodo preprint. DOI: 10.5281/zenodo.19244253.

[14] Reinhold, A. (2026). Circular associative memory architecture: A framework for emotionally-keyed AI memory systems. Zenodo. DOI: 10.5281/zenodo.19051834.

[15] NASA. (2024). First Mars crew completes yearlong simulated Red Planet NASA mission. NASA Press Release.

[16] NASA. (2026). About CHAPEA. NASA Human Research Program.

[17] Packer, C. et al. (2023). MemGPT: Towards LLMs as operating systems. arXiv preprint arXiv:2310.08560.

[18] Zhang, G. et al. (2025). Memory in the age of AI agents: A survey. arXiv preprint arXiv:2512.13564.

Appendix A: Series Overview

This paper is part of the **Applied Persistent Memory Series**, published by Lorien’s Library LLC:

Paper 1: Persistent Memory as Mission-Critical Infrastructure for Long-Duration Spaceflight (DOI: 10.5281/zenodo.19257809)

Paper 2 (this paper): Memory-Aware AI Systems for Permanent Lunar and Martian Habitation

Paper 3: Provenance-Aware Memory Architecture for Chronic Healthcare Continuity

Paper 4: Haven: Music-Mediated Emotional Continuity for Veteran Reintegration

Open-source implementation: https://github.com/LoriensLibrary/cama

License: Creative Commons Attribution 4.0 International (CC BY 4.0)